

NGINX Assignment

Title:

Set Up an HA Reverse Proxy with Load Balancing and Web Hosting

Overview


This assignment simulates a production-like High Availability (HA) environment using NGINX and Keepalived. It demonstrates how to:

- Deploy backend web servers serving different content.
- Deploy frontend load balancers using NGINX reverse proxy and round-robin load balancing.
- Implement high availability using Keepalived with VRRP and a floating Virtual IP (VIP).

Assumption: All nodes are on the same subnet and can reach each other.

Node Details

Setup 3 VM with Ubuntu 22.04 Its OS-



1- 1 vm we will use for 2 web servers

2- 2 vm we use for separate 2 load balancers i.e LB1 and LB2.

Backend Web Servers on VM 1

We are using Podman containers to run two separate nginx containers. And using the host network to access it.

1- container command-

```
podman run -d --name node-a-nginx --network host -v  
~/node-a/html:/usr/share/nginx/html:ro -v  
~/node-a/conf/nginx.conf:/etc/nginx/nginx.conf:ro docker.io/library/nginx:alpine
```

Web Server 1-

Nginx.conf

=====

```
events {}
```

```
http {  
    server {
```

```

        listen 8081;
        location / {
            root /usr/share/nginx/html;
            index index.html;
        }
    }
}
=====

```

```

c653544c0c28  docker.io/library/nginx:alpine  nginx -g daemon o...  2 days ago  Up 4 hours ago  node-a-nginx

```

2- container command-

```

podman run -d --name node-a-nginx --network host -v
~/node-a/html:/usr/share/nginx/html:ro -v
~/node-b/conf/nginx.conf:/etc/nginx/nginx.conf:ro  docker.io/library/nginx:alpine

```

Web Server 1-

Nginx.conf

```

=====
events {}

http {
    server {
        listen 8082;
        location / {
            root /usr/share/nginx/html;
            index index.html;
        }
    }
}
=====

```

Output of the both web server using curl:-

```

brijesh@brijesh:~/node-b/conf$ curl localhost:8081
<h1>Backend Server 1</h1>
brijesh@brijesh:~/node-b/conf$ curl localhost:8082
<h1>Backend Server 2</h1>
brijesh@brijesh:~/node-b/conf$

```

- **Node A (Container)** - IP: 192.168.122.240:8081
 - Serves: "Beckend Server 1"
- **Node B (Container)** - IP: 192.168.122.240:8082
 - Serves: "Beckend Server 2"

Setup Loadbalancer using Nginx-

Load Balancer Nodes (LB Tier)

- **LB1** - IP: 192.168.122.17
- **LB2** - IP: 192.168.122.137
- **VIP** - Floating IP assigned by Keepalived: 192.168.122.100

LB-1 – On Node 1

Install Nginx-

```
=====
```

```
sudo apt install -y nginx
```

```
=====
```

Default.conf

```
=====
```

```
upstream backend {
```

```
    server 192.168.122.240:8081;
```

```
    server 192.168.122.240:8082;
```

```
}
```

```
server {
```

```
    listen 80;
```

```

location / {

    proxy_pass http://backend;

    proxy_set_header Host $host;

    proxy_set_header X-Real-IP $remote_addr;

    proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;

    proxy_set_header X-Forwarded-Proto $scheme;

    proxy_next_upstream error timeout invalid_header http_500 http_502 http_503
http_504;

}

}

```

=====

```

brijesh@brijesh:/etc/nginx/conf.d$ systemctl status nginx.service
● nginx.service - A high performance web server and a reverse proxy server
   Loaded: loaded (/lib/systemd/system/nginx.service; enabled; vendor preset: enabled)
   Active: active (running) since Thu 2025-07-24 18:07:41 IST; 4h 43min ago
     Docs: man:nginx(8)
  Process: 679 ExecStartPre=/usr/sbin/nginx -t -q -g daemon on; master_process on; (code=exited, status=0/SUCCESS)
  Process: 733 ExecStart=/usr/sbin/nginx -g daemon on; master_process on; (code=exited, status=0/SUCCESS)
 Main PID: 738 (nginx)
    Tasks: 3 (limit: 4602)
   Memory: 9.1M
      CPU: 84ms
   CGroup: /system.slice/nginx.service
           └─738 "nginx: master process /usr/sbin/nginx -g daemon on; master_process on;"
             └─741 "nginx: worker process"
               └─742 "nginx: worker process"

Jul 24 18:07:41 brijesh systemd[1]: Starting A high performance web server and a reverse proxy server...
Jul 24 18:07:41 brijesh systemd[1]: Started A high performance web server and a reverse proxy server.

```

LB-2 – On Node 2

Install Nginx-

=====

```
sudo apt install -y nginx
```

=====

Default.conf

=====

upstream backend {

server 192.168.122.240:8081;

server 192.168.122.240:8082;

}

server {

listen 80;

location / {

proxy_pass http://backend;

proxy_set_header Host \$host;

proxy_set_header X-Real-IP \$remote_addr;

proxy_set_header X-Forwarded-For \$proxy_add_x_forwarded_for;

proxy_set_header X-Forwarded-Proto \$scheme;

proxy_next_upstream error timeout invalid_header http_500 http_502 http_503
http_504;

}

}

[illegible]

Note- Remove LB1 and LB2 nginx static web pages by using below command and then reload Nginx to apply changes-

=====

```
sudo rm /etc/nginx/sites-enabled/default
```

=====

```
sudo systemctl reload nginx
```

=====

Keepalived Setup on the both LB1 and LB2 Node:-

3. Configure Keepalived (High Availability)

Installation-

=====

```
sudo apt install keepalived -y
```

=====

LB1 keepalive we use as Master-

LB1 - /etc/keepalived/keepalived.conf

=====

```
vrrp_instance VI_1 {  
    state MASTER  
  
    interface enp1s0  
  
    virtual_router_id 51  
  
    priority 150  
  
    advert_int 1  
  
    authentication {  
        auth_type PASS  
        auth_pass secret123  
    }  
  
    virtual_ipaddress {  
        192.168.122.100  
    }  
}
```

=====

```
sudo systemctl enable --now keepalived
```

=====

LB2 we use as Backup-

LB2 - /etc/keepalived/keepalived.conf

=====

```
vrp_instance VI_1 {  
    state BACKUP  
  
    interface enp1s0  
  
    virtual_router_id 51  
  
    priority 100  
  
    advert_int 1  
  
    authentication {  
        auth_type PASS  
        auth_pass boss123  
    }  
  
    virtual_ipaddress {  
        192.168.122.100  
    }  
}
```

=====

```
sudo systemctl enable --now keepalived
```

=====

Testing Instructions


1. Test VIP & Load Balancing

From any host or browser:

```
curl http://192.168.122.100
```


Repeat the curl command multiple times. You should see alternating outputs:

- "Backend Server 1"
- "Backend Server 2"



Backend Server 1

On Reload-



Backend Server 2

2. Simulate Failover

On LB1:

1- ip a

```
Jul 24 23:52
brijesh-nginx-lb1 on QEMU/KVM

s Terminal Jul 24 23:52
brijesh@brijesh: ~

brijesh@brijesh:~$ ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: enp1s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 52:54:00:4b:f4:40 brd ff:ff:ff:ff:ff:ff
    inet 192.168.122.137/24 brd 192.168.122.255 scope global dynamic noprefixroute enp1s0
        valid_lft 3570sec preferred_lft 3570sec
    inet 192.168.122.100/32 scope global enp1s0
        valid_lft forever preferred_lft forever
    inet6 fe80::4bac:2edd:48a8:ad43/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
brijesh@brijesh:~$
```

sudo systemctl stop keepalived

- Run ip a on LB2 to verify VIP 192.168.122.100 has been assigned.

```
Jul 24 23:51
brijesh-nginx-lb2 on QEMU/KVM

Terminal Jul 24 23:51

brijesh@brijesh:~$ ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: enp1s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 52:54:00:c3:98:8f brd ff:ff:ff:ff:ff:ff
    inet 192.168.122.17/24 brd 192.168.122.255 scope global dynamic noprefixroute
        valid_lft 2340sec preferred_lft 2340sec
    inet 192.168.122.100/32 scope global enp1s0
        valid_lft forever preferred_lft forever
    inet6 fe80::fcec:65f4:803a:7ce3/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
brijesh@brijesh:~$
```

- VIP traffic will now be handled by LB2.

3. Recovery

On LB1:

sudo systemctl start keepalived

- VIP will automatically fail back to LB1 due to higher priority.

Output

Round-robin Load Balancing

curl http://192.168.122.100

Backend Server 1

curl http://192.168.122.100

Backend server 2

VIP Before & After Failover

- Before failover:
 - LB1 has VIP 192.168.122.100
 - After stopping keepalived on LB1:
 - LB2 takes over VIP
-

Conclusion

This assignment demonstrates a highly available web architecture using:

- NGINX as a reverse proxy and load balancer
- Podman containers as backend nodes
- Keepalived for failover using VRRP and VIP

All components work together to simulate real-world redundancy and load balancing in a multi-tier deployment.